

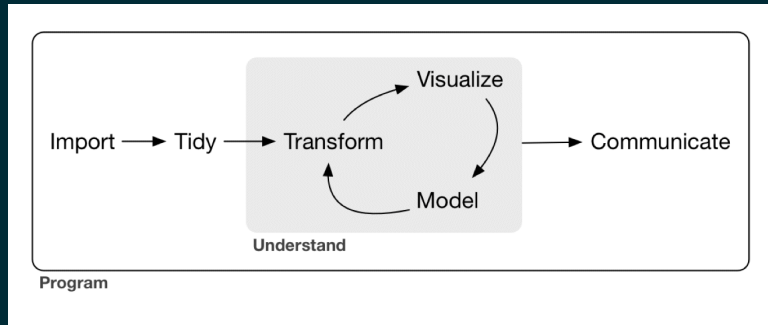
INTRODUCTION TO DATA SCIENCE WITH R

Session 6b: Tidying data

Ina Bornkessel-Schlesewsky

January 20, 2023

RECALL



- **import** data (into R)
- **tidy** data
 - bring it into a consistent format that can be used for multiple purposes (each column = variable; each row = observation)
 - lets you focus on understanding the data rather than which format you need
- **transform** data
 - e.g. focus on observations of interest (such as those from a particular location), create new variables (such as speed from distance and time), compute summary statistics
- **visualise** data
 - essential for understanding
- **model** data
 - use (statistical) models to answer your questions about the data
- **communicate** insights

DATA STRUCTURE

- Most datasets are made up of rows and columns
- Many ways to structure the same data (see figure)
- Datasets are collections of values (either numbers or strings); these belong to a *variable* and an *observation*
- “A variable contains all values that measure the same underlying attribute (like height, temperature, duration) across units.”
- “An observation contains all values measured on the same unit (like a person, or a day, or a race) across attributes.” (Wickham, 2014)

	treatmenta	treatmentb
John Smith	—	2
Jane Doe	16	11
Mary Johnson	3	1

Table 1: Typical presentation dataset.

	John Smith	Jane Doe	Mary Johnson
treatmenta	—	16	3
treatmentb	2	11	1

Table 2: The same data as in Table 1 but structured differently.

figure from Wickham (2014)

Wickham, H. (2014). Tidy Data. Journal of Statistical Software, 59(10), 1 - 23.
doi:<http://dx.doi.org/10.18637/jss.v059.i10>

DATA STRUCTURE

- The same dataset but with observations in rows and variables in columns

person	treatment	result
John Smith	a	—
Jane Doe	a	16
Mary Johnson	a	3
John Smith	b	2
Jane Doe	b	11
Mary Johnson	b	1

figure from Wickham (2014)

TIDY DATA

CHARACTERISTICS OF TIDY DATA:*

- Each variable has its own column
- Each observation has its own row
- Each value has its own cell

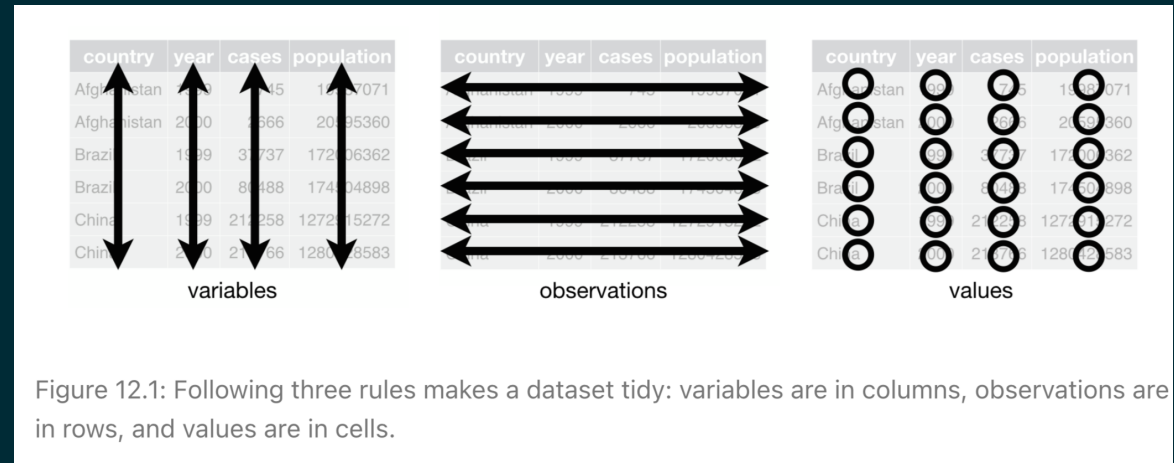


figure from R4DS

DIFFERENT DATA FORMATS

- The three data frames on the next slide show example data provided in the `tidyr` package, which is part of the `tidyverse`.

“all display the number of TB cases documented by the World Health Organization in Afghanistan, Brazil, and China between 1999 and 2000”

(from `?tidyr::table1`)

- Which of these tables is tidy?

If you want to inspect the data yourself, you can access them via `table1`, `table2` and `table3`.

DIFFERENT DATA FORMATS

```
[1] "table1"
```

```
# A tibble: 6 × 4
```

	country	year	cases	population
	<chr>	<int>	<int>	<int>
1	Afghanistan	1999	745	19987071
2	Afghanistan	2000	2666	20595360
3	Brazil	1999	37737	172006362
4	Brazil	2000	80488	174504898
5	China	1999	212258	1272915272
6	China	2000	213766	1280428583

```
[1] "table2"
```

```
# A tibble: 6 × 4
```

	country	year	type	count
	<chr>	<int>	<chr>	<int>
1	Afghanistan	1999	cases	745
2	Afghanistan	1999	population	19987071
3	Afghanistan	2000	cases	2666
4	Afghanistan	2000	population	20595360
5	Brazil	1999	cases	37737
6	Brazil	1999	population	172006362

```
[1] "table3"
```

```
# A tibble: 6 × 3
```

	country	year	rate
*	<chr>	<int>	<chr>
1	Afghanistan	1999	745/19987071
2	Afghanistan	2000	2666/20595360
3	Brazil	1999	37737/172006362
4	Brazil	2000	80488/174504898
5	China	1999	212258/1272915272
6	China	2000	213766/1280428583

DIFFERENT DATA FORMATS

- If you said `table1`, you were right!
- Another possibility is for data to be spread across two data frames
- See the next slide for an example (`table4a` and `table4b` from `tidyr`)

DIFFERENT DATA FORMATS

```
[1] "table4a"
# A tibble: 3 × 3
  country    `1999`    `2000`
*   <chr>      <int>    <int>
1 Afghanistan      745      2666
2 Brazil          37737    80488
3 China          212258   213766
```

```
[1] "table4b"
# A tibble: 3 × 3
  country    `1999`    `2000`
*   <chr>      <int>    <int>
1 Afghanistan 19987071  20595360
2 Brazil      172006362 174504898
3 China      1272915272 1280428583
```

ADVANTAGES TO WORKING WITH TIDY DATA

1. Having one consistent format for data makes it easier to learn the tools required for analysis (which can have a certain uniformity). The **tidyverse** packages, for example, are designed to work with tidy data (who would have thought! 🤔)
2. It is advantageous for variables to be placed in columns because this caters to R's vectorised nature. (Most R-functions work with vectors of values.)

DATA ARE OFTEN UNTIDY

COMMON PROBLEMS

1. Variables are spread across multiple columns
2. Observations are spread across multiple rows

THE SOLUTION

- functions `pivot_longer()` and `pivot_wider()` in `tidyr`!

Note: this doesn't mean that non-tidy data are "bad". There can be many reasons for why a dataset is in a non-tidy format, e.g. ease of data entry if this is being done manually.

PIVOT TO LONGER

Common problem: column names are values of a variable rather than variables

Example: `table4a`

```
# A tibble: 3 × 3
  country    `1999` `2000`
  <chr>      <int> <int>
1 Afghanistan    745   2666
2 Brazil        37737  80488
3 China         212258 213766
```

Solution: *pivot* these columns to new variables, rendering the dataset longer

We need:

- the columns with values as names (**1999** and **2000**)
- the name of the variable to move the column names to (**year**)
- the name of the variable to move the column values to (**cases**)

PIVOT TO LONGER

```
table4a |>
  pivot_longer(c(`1999`, `2000`), names_to = "year", values_to = "cases")
```

```
# A tibble: 6 × 3
  country    year  cases
  <chr>      <chr> <int>
1 Afghanistan 1999     745
2 Afghanistan 2000    2666
3 Brazil      1999   37737
4 Brazil      2000   80488
5 China       1999  212258
6 China       2000  213766
```

PIVOT TO LONGER

country	year	cases
Afghanistan	1999	745
Afghanistan	2000	2666
Brazil	1999	37737
Brazil	2000	80488
China	1999	212258
China	2000	213766

country	1999	2000
Afghanistan	745	2666
Brazil	37737	80488
China	212258	213766

table4

Figure 12.2: Pivoting `table4` into a longer, tidy form.

from R4DS

Exercise: try doing the same thing with `table4b`!

EXCURSUS 1: JOINING TABLES

We can easily join the longer versions of `table4a` and `table4b` using `left_join()` (more on joining operations later):

```
tidy4a <- table4a %>%  
  pivot_longer(c(`1999`, `2000`),  
               names_to = "year",  
               values_to = "cases")  
  
tidy4b <- table4b %>%  
  pivot_longer(c(`1999`, `2000`),  
               names_to = "year",  
               values_to = "population")  
  
left_join(tidy4a, tidy4b)
```

```
# A tibble: 6 × 4  
  country    year  cases population  
  <chr>    <chr> <int>      <int>  
1 Afghanistan 1999     745   19987071  
2 Afghanistan 2000    2666   20595360  
3 Brazil      1999   37737   172006362  
4 Brazil      2000   80488   174504898  
5 China       1999  212258  1272915272  
6 China       2000  213766  1280428583
```

EXCURSUS 2: ADDITIONAL TOOLS FOR CLEANING DATA

- the `{janitor}` package includes a number of useful functions for cleaning data
- one of these is `clean_names()`, which cleans up problematic variable names (e.g. names with spaces, starting with a digit etc.)

```
table4a |>
  janitor::clean_names()

# A tibble: 3 × 3
  country    x1999    x2000
* <chr>      <int>    <int>
1 Afghanistan     745     2666
2 Brazil        37737    80488
3 China         212258   213766
```


PIVOT TO WIDER

- `pivot_wider()` is the counterpart of `pivot_longer()` which you need when observations are spread across multiple rows such as in `table2`
- here, the table needs to be made wider

```
# A tibble: 6 × 4
  country    year type      count
  <chr>    <int> <chr>    <int>
1 Afghanistan 1999 cases       745
2 Afghanistan 1999 population 19987071
3 Afghanistan 2000 cases       2666
4 Afghanistan 2000 population 20595360
5 Brazil      1999 cases       37737
6 Brazil      1999 population 172006362
```

PIVOT TO WIDER

To tidy `table2` we need:

- the column to take variables from (`type`)
- the column to take values from (`count`)

```
table2 |>
  pivot_wider(names_from = type, values_from = count)
```

```
# A tibble: 6 × 4
  country    year cases population
  <chr>      <int> <int>      <int>
1 Afghanistan 1999     745  19987071
2 Afghanistan 2000    2666  20595360
3 Brazil      1999   37737  172006362
4 Brazil      2000   80488  174504898
5 China       1999  212258  1272915272
6 China       2000  213766  1280428583
```

PIVOT TO WIDER

country	year	cases
Afghanistan	1999	745
Afghanistan	2000	2666
Brazil	1999	37737
Brazil	2000	80488
China	1999	212258
China	2000	213766

country	1999	2000
Afghanistan	745	2666
Brazil	37737	80488
China	212258	213766

table4

Figure 12.2: Pivoting `table4` into a longer, tidy form.

from R4DS

WHAT'S UP WITH TABLE3?

```
# A tibble: 6 × 3
  country    year rate
* <chr>    <int> <chr>
1 Afghanistan 1999 745/19987071
2 Afghanistan 2000 2666/20595360
3 Brazil      1999 37737/172006362
4 Brazil      2000 80488/174504898
5 China       1999 212258/1272915272
6 China       2000 213766/1280428583
```

MULTIPLE VARIABLES IN ONE COLUMN

- in `table3`, the `rate` column contains both cases and population
- to deal with this problem, we can use the `separate()` function
- it allows us to easily split a column according to a delimiting character (here, the `"/"`)
- note how `separate` is clever enough to correctly guess the delimiting character – it looks for a non-alphanumeric character by default (to specify it manually, use `sep = "/"`)

```
table3 |>
  separate(rate, into = c("cases", "population"))
```

```
# A tibble: 6 × 4
  country    year cases population
  <chr>    <int> <chr>    <chr>
1 Afghanistan 1999  745    19987071
2 Afghanistan 2000 2666    20595360
3 Brazil      1999 37737   172006362
4 Brazil      2000 80488   174504898
5 China       1999 212258  1272915272
6 China       2000 213766  1280428583
```

MULTIPLE VARIABLES IN ONE COLUMN

- by default, **separate** retains the original column type (character in this case)
- we can ask it to try to convert to a more suitable type using the **convert** parameter

```
table3 |>  
  separate(rate, into = c("cases", "population"), convert = TRUE)
```

```
# A tibble: 6 × 4  
  country    year cases population  
  <chr>    <int> <int>      <int>  
1 Afghanistan 1999     745 19987071  
2 Afghanistan 2000    2666 20595360  
3 Brazil      1999   37737 172006362  
4 Brazil      2000   80488 174504898  
5 China       1999  212258 1272915272  
6 China       2000  213766 1280428583
```

RESOURCES

- Chapter 6: Data tidying of R4DS
- Chapter 21: Joins of R4DS provides more information on joining data frames, which is beyond the scope of this course

